County level population by race ethicity 2010-2019

Name: SREEPRATHAP M

Organization: Entri Elivate

Overview of Problem Statement

The population by race and ethicity census data of 2010-2019 of 'US' is a collection of population of different ethnic groups like the white population, Asian population, Native American population, and so on. By predicting the major ethnic group or race population, we can identify which population of people contributes more to the country in various aspects.

Objective

To figure out major ethnic group predictions using machine learning techniques.

Data Description

Source:

County level population by race ethnicity 2010-2019

Features:

(1).fips:- State and county FIPS code,
(2).stfips:- State FIPS code,
(3).cofips:- County FIPS code,
(4).state_abbrev:- Stae abbrevation,
(5).state:- State name,
(6).county:- County name,
(7).year:- Year,
(8).pop:- Total population(all races),
(9).white_pop:- White population,
(10).black_pop:- Black population,
(11).asian_pop:- Asian population,
(12).indian_pop:- Naive American or Alaska native population,
(13).pacific_pop:- Native Hawaiian or other Pacific islander Population,
(14).two_pop:- Tow or more races population,
(15).not_hisp_pop:- Non-Hispanic population (independent of race),
(16).hisp_pop:- Hispanic population(independent of race)

Data Collection

```
In [6]:   # importing library
          import numpy as np
          import pandas as pd
```

```
In [7]:   # Reading downloaded Dataset from loacl directory
          data = pd.read_csv('US_county_census_est_race_eth_2010_2019.csv')
          # Converging Data to panda DataFrame
          df = pd.DataFrame(data)
```

```
In [8]:   ## For result and title printing, create a custom definition.
          def print_title(title):
              print(f'\n{'-'*60}\n\033[1m{title}\033[0m')
          def print_section(title):
              print(f'{'-'*60}\n{title}\n{'-'*60}')
```

```
In [9]:   df.head(3)
```

Out[9]:

	FIPS	STFIPS	COFIPS	state_abbrev	state	county	year	pop	white_pop	black_pop	asian_pop	indian
0	1001	1	1	AL	Alabama	Autauga	2010	54571	43297	9689	484	
1	1001	1	1	AL	Alabama	Autauga	2011	55227	43699	9883	514	
2	1001	1	1	AL	Alabama	Autauga	2012	54954	43315	9949	552	

```
In [10]:  df.tail(3)
```

Out[10]:

	FIPS	STFIPS	COFIPS	state_abbrev	state	county	year	pop	white_pop	black_pop	asian_pop	in
31407	56045	56	45	WY	Wyoming	Weston	2017	6968	6558	44	97	
31408	56045	56	45	WY	Wyoming	Weston	2018	6924	6474	47	109	
31409	56045	56	45	WY	Wyoming	Weston	2019	6927	6454	48	117	

Data Preprocessing - Data Cleaing

```
In [12]:  df1 = pd.DataFrame(df)
          print_section(f'Since column "state_abbrev" and "state" columns are same, \ndorping "sta
          df1 = df1.drop("state_abbrev", axis=1)
          print_section("DataFrame after dropping the column")
          df1.head(3)
```

```
          ------------------------------------------------------------
          Since column "state_abbrev" and "state" columns are same,
          dorping "state_abbrev" columns and creating df1 the from datafrme
          ------------------------------------------------------------
          ------------------------------------------------------------
          DataFrame after dropping the column
          ------------------------------------------------------------
```

Out[12]:

	FIPS	STFIPS	COFIPS	state	county	year	pop	white_pop	black_pop	asian_pop	indian_pop	pacific_
0	1001	1	1	Alabama	Autauga	2010	54571	43297	9689	484	258	
1	1001	1	1	Alabama	Autauga	2011	55227	43699	9883	514	261	
2	1001	1	1	Alabama	Autauga	2012	54954	43315	9949	552	275	

```
In [13]:  # data frme information
```

```
print_title("DataFrame Information")
print_section(df1.info())
```

```
------------------------------------------------------------
DataFrame Information
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 31410 entries, 0 to 31409
Data columns (total 15 columns):
 #   Column         Non-Null Count  Dtype
---  ------         --------------  -----
 0   FIPS           31410 non-null  int64
 1   STFIPS         31410 non-null  int64
 2   COFIPS         31410 non-null  int64
 3   state          31410 non-null  object
 4   county         31410 non-null  object
 5   year           31410 non-null  int64
 6   pop            31410 non-null  int64
 7   white_pop      31410 non-null  int64
 8   black_pop      31410 non-null  int64
 9   asian_pop      31410 non-null  int64
 10  indian_pop     31410 non-null  int64
 11  pacific_pop    31410 non-null  int64
 12  two_pop        31410 non-null  int64
 13  not_hisp_pop   31410 non-null  int64
 14  hisp_pop       31410 non-null  int64
dtypes: int64(13), object(2)
memory usage: 3.6+ MB
------------------------------------------------------------
None
------------------------------------------------------------
```

In [14]:
```
# Data frame Description
print_title("DataFrame Description")
print_section(df1.describe())
```

```
------------------------------------------------------------
DataFrame Description
------------------------------------------------------------
                FIPS        STFIPS        COFIPS          year           pop  \
count   31410.000000  31410.000000  31410.000000  31410.000000  3.141000e+04
mean    30389.820121     30.286215    103.605540   2014.500000  1.014097e+05
std     15158.803727     15.140671    107.690218      2.872327  3.251245e+05
min      1001.000000      1.000000      1.000000   2010.000000  8.200000e+01
25%     18179.000000     18.000000     35.000000   2012.000000  1.098500e+04
50%     29177.000000     29.000000     79.000000   2014.500000  2.573350e+04
75%     45081.000000     45.000000    133.000000   2017.000000  6.741675e+04
max     56045.000000     56.000000    840.000000   2019.000000  1.010571e+07

           white_pop     black_pop     asian_pop     indian_pop    pacific_pop  \
count   3.141000e+04  3.141000e+04  3.141000e+04   31410.000000   31410.000000
mean    7.844218e+04  1.335439e+04  5.543686e+03    1264.135371     236.176313
std     2.333952e+05  5.778493e+04  4.089464e+04    5203.210936    2150.884073
min     2.400000e+01  0.000000e+00  0.000000e+00       0.000000       0.000000
25%     9.105000e+03  1.170000e+02  4.700000e+01      64.000000       4.000000
50%     2.217200e+04  8.400000e+02  1.560000e+02     179.000000      14.000000
75%     5.846450e+04  5.697750e+03  7.750000e+02     618.750000      60.000000
max     7.181207e+06  1.311698e+06  1.545445e+06  146005.000000   95285.000000

             two_pop  not_hisp_pop      hisp_pop
count   31410.000000  3.141000e+04  3.141000e+04
mean     2569.121999  8.370397e+04  1.770571e+04
std     10265.356718  2.222491e+05  1.228180e+05
min         0.000000  6.400000e+01  0.000000e+00
25%       156.000000  9.944500e+03  3.240000e+02
50%       392.000000  2.386300e+04  1.015500e+03
75%      1348.750000  6.290925e+04  4.764500e+03
```

```
       max       315568.000000   5.211947e+06   4.899383e+06
       ----------------------------------------------------------------
```

In [15]:
```python
# Finding null value for each features
print_title("Null values in DataFrame")
print_section(df1.isnull().sum())
```

```
----------------------------------------------------------------
Null values in DataFrame
----------------------------------------------------------------
FIPS              0
STFIPS            0
COFIPS            0
state             0
county            0
year              0
pop               0
white_pop         0
black_pop         0
asian_pop         0
indian_pop        0
pacific_pop       0
two_pop           0
not_hisp_pop      0
hisp_pop          0
dtype: int64
----------------------------------------------------------------
```

In [16]:
```python
# Features data type
print_title("Data Types of Dataframe Varible columns")
print_section(df1.dtypes)
```

```
----------------------------------------------------------------
Data Types of Dataframe Varible columns
----------------------------------------------------------------
FIPS             int64
STFIPS           int64
COFIPS           int64
state           object
county          object
year             int64
pop              int64
white_pop        int64
black_pop        int64
asian_pop        int64
indian_pop       int64
pacific_pop      int64
two_pop          int64
not_hisp_pop     int64
hisp_pop         int64
dtype: object
----------------------------------------------------------------
```

In [17]:
```python
# Finding Duplicated valus
print_section(f"Duplicated values: {df1.duplicated().sum()}")
```

```
----------------------------------------------------------------
Duplicated values: 0
----------------------------------------------------------------
```

Outlier Removel

In [19]:
```python
# Creating custom definition to remove outliers using IQR method
def outliers(data):
    for col in data.select_dtypes(include=['int64','float64']).columns:
        Q1 = data[col].quantile(0.25)
```

```
        Q3 = data[col].quantile(0.75)
        IQR = Q3 - Q1

        lower = Q1 - (1.5*IQR)
        upper = Q3 + (1.5*IQR)

        # Capping
        data[col] = data[col].apply(lambda x: lower if x < lower else upper if x > upper

    return data
```

In [20]:
```
df2 = pd.DataFrame(df1) # data frame after outlier remove
outliers(df2)
```

Out[20]:

	FIPS	STFIPS	COFIPS	state	county	year	pop	white_pop	black_pop	asian_pop	indian_pop
0	1001	1	1.0	Alabama	Autauga	2010	54571.0	43297.0	9689.0	484.0	258.0
1	1001	1	1.0	Alabama	Autauga	2011	55227.0	43699.0	9883.0	514.0	261.0
2	1001	1	1.0	Alabama	Autauga	2012	54954.0	43315.0	9949.0	552.0	275.0
3	1001	1	1.0	Alabama	Autauga	2013	54727.0	42943.0	9984.0	561.0	279.0
4	1001	1	1.0	Alabama	Autauga	2014	54893.0	42945.0	10103.0	573.0	279.0
...	...	...	...	...	...	...	...	...	...	...	...
31405	56045	56	45.0	Wyoming	Weston	2015	7208.0	6835.0	39.0	81.0	107.0
31406	56045	56	45.0	Wyoming	Weston	2016	7220.0	6826.0	38.0	88.0	108.0
31407	56045	56	45.0	Wyoming	Weston	2017	6968.0	6558.0	44.0	97.0	114.0
31408	56045	56	45.0	Wyoming	Weston	2018	6924.0	6474.0	47.0	109.0	125.0
31409	56045	56	45.0	Wyoming	Weston	2019	6927.0	6454.0	48.0	117.0	131.0

31410 rows × 15 columns

Skewness

In [22]:
```
skewness = df2.select_dtypes(include = ['int64','float64']).skew()
print_section("\033[1mSkewness of numerical Features \033[0m")
print_section(skewness)
```

```
--------------------------------------------------------------
Skewness of numerical Features
--------------------------------------------------------------
--------------------------------------------------------------
FIPS            -0.080233
STFIPS          -0.082250
COFIPS           0.968919
year             0.000000
pop              1.148729
white_pop        1.164347
black_pop        1.221158
asian_pop        1.203443
indian_pop       1.158144
pacific_pop      1.198770
two_pop          1.190157
not_hisp_pop     1.168889
hisp_pop         1.208730
dtype: float64
--------------------------------------------------------------
```

Exploratory Data analysis (EDA)

```
In [24]: # Importing the plotting library
         import matplotlib.pyplot as plt
         import seaborn as sns
```

```
In [25]: plt.figure(figsize=(8,5))
         sns.boxplot(data=df2)
         plt.title(f'''Boxplot of Numarical features
         outlier removed
         ''')
         plt.xticks(rotation = 90)
         plt.show()
```


Boxplot of Numarical features
outlier removed

```
In [26]: ax=sns.histplot(df2,kde=True,linewidth=0,legend=True)
         plt.title('Histplot of dataframe')
         sns.move_legend(ax, "upper left", bbox_to_anchor=(1, 1))
         plt.show()
```

Histplot of dataframe
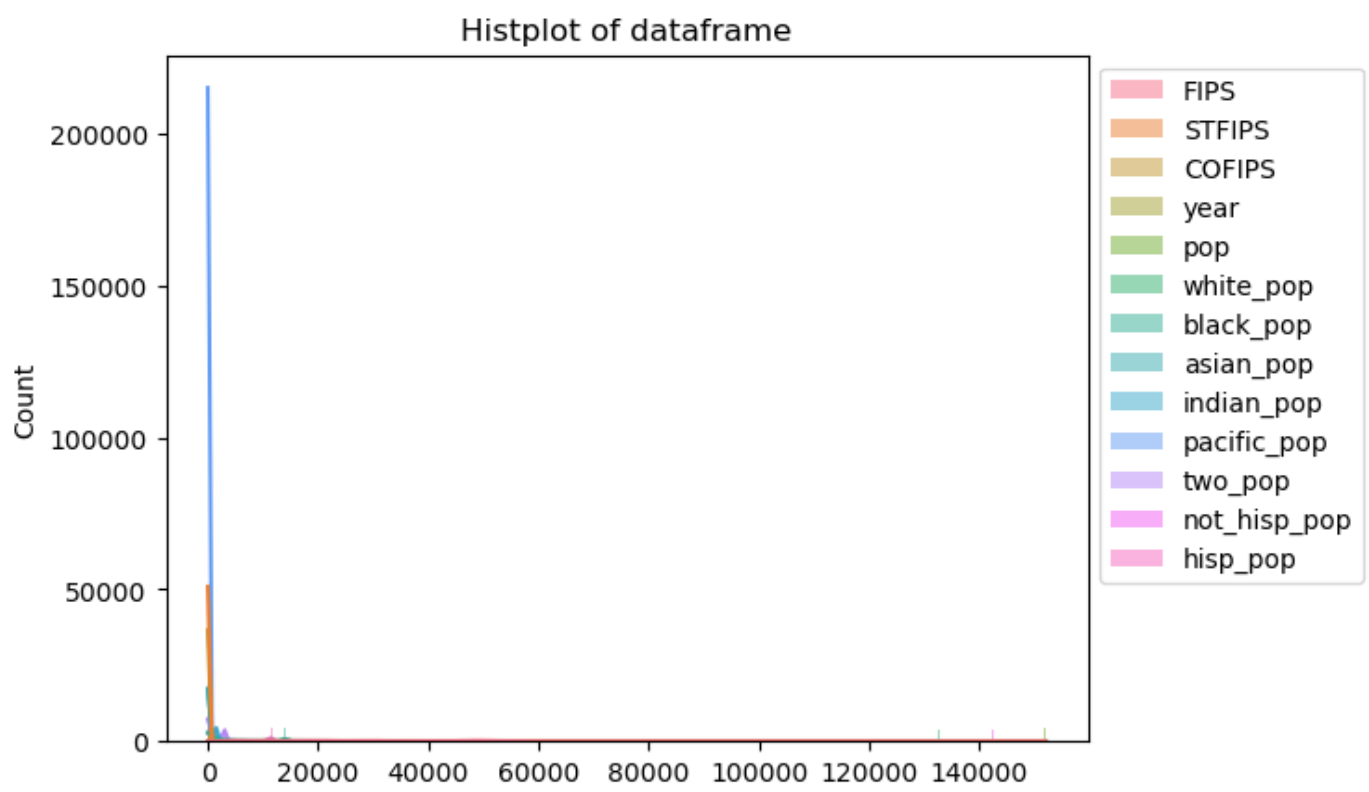
```
In [27]:  #correlation of data frame
          cor_df = df2.select_dtypes(include = ['int64','float64']).corr()
          #correlation plot
          plt.figure(figsize=(12,10))
          sns.heatmap(cor_df,annot=True,annot_kws={'size': 7})
          plt.title('Correlation plot population dataframe')
          plt.show()
```

Correlation plot population dataframe

```
# Barchard for data visualization
nu_col=df2.select_dtypes(include=['int64','float64'])
df3 = pd.DataFrame(df2.select_dtypes(include=['int64','float64']))
sns.barplot(data=df3)
plt.xticks(rotation = 90)
plt.title('Bar chart of population dataframe')
plt.show()
```

Bar chart of population dataframe


```
In [29]:   # ECDF cahrt for data visualization
           ax=sns.ecdfplot(data=df3,legend=True)
           sns.move_legend(ax, "upper left", bbox_to_anchor=(1, 1))
           plt.xticks(rotation = 90)
           plt.title('ECDF chart of population dataframe')
           plt.show()
```

ECDF chart of population dataframe

Feature Engineering

Encoding

Encoding of object-type columns

```
In [31]:   # importing category encoder from library
           from category_encoders import TargetEncoder
```

```
In [32]:   df4 = pd.DataFrame(df2)
           ## finding Object type columns and count the number of unique values
           col_object_type=df4.select_dtypes(include=['object']).columns.tolist()
           print_title('Object type columns in DataFrame')
           print_section(pd.DataFrame({'Colums':col_object_type}))
           print_title('Counted unique values in object type columns')
           for column_name in col_object_type:
             print(f'{column_name} has {df4[column_name].nunique()} unique values')
```

```
-----------------------------------------------------------
Object type columns in DataFrame
-----------------------------------------------------------
    Colums
0    state
1   county
-----------------------------------------------------------


-----------------------------------------------------------
Counted unique values in object type columns
state has 50 unique values
county has 1876 unique values
```

```
In [33]:   print_title('Aggrigation of Cont and Mean of object type column to target column')
```

```
for column_name in col_object_type:
    print_section(df4['pop'].groupby(df4[column_name]).agg(['count','mean']))
```

Aggrigation of Cont and Mean of object type column to target column
----

	count	mean
state		
Alabama	670	54314.985821
Alaska	290	20300.575000
Arizona	150	101803.436667
Arkansas	750	34223.004333
California	580	107221.903879
Colorado	640	39205.986328
Connecticut	80	147631.015625
Delaware	30	152064.375000
Florida	670	94364.632090
Georgia	1590	42915.739072
Hawaii	50	105304.565000
Idaho	440	29997.327273
Illinois	1020	47549.117402
Indiana	920	52840.632065
Iowa	990	27400.281061
Kansas	1050	19803.138095
Kentucky	1200	30299.577708
Louisiana	640	54741.042188
Maine	160	71594.773438
Maryland	240	101035.480729
Massachusetts	140	124844.246429
Michigan	830	61218.862952
Minnesota	870	39366.908764
Mississippi	820	34445.540701
Missouri	1150	34649.018261
Montana	560	18284.945759
Nebraska	930	14168.279839
Nevada	170	37332.132353
New Hampshire	100	92950.175000
New Jersey	210	141338.985119
New Mexico	330	45511.919697
New York	620	95734.444355
North Carolina	1000	69040.397125
North Dakota	530	13482.289387
Ohio	880	75729.046307
Oklahoma	770	34584.166558
Oregon	360	62544.022917
Pennsylvania	670	92346.280224
Rhode Island	50	112496.330000
South Carolina	460	73805.552446
South Dakota	660	12452.670833
Tennessee	950	47494.407105
Texas	2540	40466.837598
Utah	290	45889.747845
Vermont	140	44047.512500
Virginia	1330	43871.389850
Washington	390	72563.562821
West Virginia	550	32728.988636
Wisconsin	720	57258.012153
Wyoming	230	25117.400000

----

	count	mean
county		
Abbeville	10	24834.5000
Acadia Parish	10	62288.4000
Accomack	10	32896.4000
Ada	10	152064.3750

```
Adair                          40    18568.0250
...                           ...           ...
Yukon-Koyukuk Census Area      10     5474.7000
Yuma                           20    81062.6875
Zapata                         10    14280.3000
Zavala                         10    12020.1000
Ziebach                        10     2818.4000

[1876 rows x 2 columns]
```

In [34]:
```python
print_section(f'Since high cardinality in state and county coluns Target\nEncoding is mo
encoder=TargetEncoder()
```

```
------------------------------------------------------------
Since high cardinality in state and county coluns Target
Encoding is moste prefered
------------------------------------------------------------
```

In [35]:
```python
df5 = pd.DataFrame(df4)
df4.head(3)
```

Out[35]:

	FIPS	STFIPS	COFIPS	state	county	year	pop	white_pop	black_pop	asian_pop	indian_pop	pacific
0	1001	1	1.0	Alabama	Autauga	2010	54571.0	43297.0	9689.0	484.0	258.0	
1	1001	1	1.0	Alabama	Autauga	2011	55227.0	43699.0	9883.0	514.0	261.0	
2	1001	1	1.0	Alabama	Autauga	2012	54954.0	43315.0	9949.0	552.0	275.0	

In [36]:
```python
print_section(f'Since high cardinality in state and county coluns Target\nEncoding is mo
encoder.fit(df4['state'],df4['pop'])
df5['state'] = encoder.transform(df4['state'],df4['pop'])
encoder.fit(df4['county'],df4['pop'])
df5['county'] = encoder.transform(df4['county'],df4['pop'])
df5.head(3)
```

```
------------------------------------------------------------
Since high cardinality in state and county coluns Target
Encoding is moste prefered
------------------------------------------------------------
```

Out[36]:

	FIPS	STFIPS	COFIPS	state	county	year	pop	white_pop	black_pop	asian_pop	indian_po
0	1001	1	1.0	54314.985821	50281.243963	2010	54571.0	43297.0	9689.0	484.0	258.
1	1001	1	1.0	54314.985821	50281.243963	2011	55227.0	43699.0	9883.0	514.0	261.
2	1001	1	1.0	54314.985821	50281.243963	2012	54954.0	43315.0	9949.0	552.0	275.

Feature Scaling

Min Max scaling

In [38]:
```python
# importing MinMaxScaler library for feature Scaling
from sklearn.preprocessing import MinMaxScaler
```

In [39]:
```python
df6 = pd.DataFrame(df5)
```

In [40]:
```python
scaling=MinMaxScaler()
numerical_col = df6.select_dtypes(include=['int64','float64'])
```

```python
df6 = scaling.fit_transform(numerical_col)
df6=pd.DataFrame(df6, columns=numerical_col.columns, index=df2.index) #datafrme after mi
```

PowerTransformer

In [42]:
```python
# importing powertransformer library
from sklearn.preprocessing import PowerTransformer
```

In [43]:
```python
pt = PowerTransformer(method='yeo-johnson',standardize=True)
numeriacal_features = df6.select_dtypes(include=['int64','float64']).columns
df6[numeriacal_features] = pt.fit_transform(df6[numeriacal_features])
```

In [44]:
```python
df6.head(5)
```

Out[44]:

	FIPS	STFIPS	COFIPS	state	county	year	pop	white_pop	black_pop	asian_pop	india
0	-1.928755	-1.92287	-1.726894	0.458593	0.274188	-1.610567	0.616111	0.527443	1.420259	0.566839	0.0
1	-1.928755	-1.92287	-1.726894	0.458593	0.274188	-1.230735	0.629712	0.537918	1.432322	0.625005	0.0
2	-1.928755	-1.92287	-1.726894	0.458593	0.274188	-0.860716	0.624072	0.527914	1.436322	0.693883	0.0
3	-1.928755	-1.92287	-1.726894	0.458593	0.274188	-0.499397	0.619361	0.518138	1.438423	0.709458	0.7
4	-1.928755	-1.92287	-1.726894	0.458593	0.274188	-0.145869	0.622808	0.518191	1.445458	0.729805	0.7

In [45]:
```python
skewness = df6.skew()
print_section("\033[1mSkewness of Features in dataframe after Scaling\033[0m")
print_section(skewness)
```

```
--------------------------------------------------------------
Skewness of Features in dataframe after Scaling
--------------------------------------------------------------
--------------------------------------------------------------
FIPS             -0.070539
STFIPS           -0.071084
COFIPS            0.098267
state             0.033800
county           -0.000077
year             -0.066401
pop               0.336701
white_pop         0.336056
black_pop         0.557582
asian_pop         0.554942
indian_pop        0.438283
pacific_pop       0.486517
two_pop           0.444147
not_hisp_pop      0.335262
hisp_pop          0.505745
dtype: float64
--------------------------------------------------------------
```

In [46]:
```python
ax=sns.histplot(df6,kde=True,linewidth=0,legend=True)
plt.title('Histplot of dataframe after MinMaxScaling and Powertransformer')
sns.move_legend(ax, "upper left", bbox_to_anchor=(1, 1))
plt.show()

ax=sns.histplot(df2,kde=True,linewidth=0,legend=True)
plt.title('Histplot of dataframe befor Preprocessing')
sns.move_legend(ax, "upper left", bbox_to_anchor=(1, 1))
plt.show()
```

Histplot of dataframe after MinMaxScaling and Powertransformer

Histplot of dataframe befor Preprocessing

In [47]:
```python
#correlation of data frame after EDA
cor_df = df6.corr()
#correlation plot
plt.figure(figsize=(12,10))
sns.heatmap(cor_df,annot=True,annot_kws={'size': 7})
plt.title('Correlation plot')
plt.show()
```

Correlation plot

Feature Selection

```
In [49]:  # importing SelectKBest library for feature selection
          from sklearn.feature_selection import SelectKBest,f_regression
```

```
In [50]:  #converting df6 to new dataset name for futher process
          df_pop = pd.DataFrame(df6)
```

```
In [185…  df_pop.head(5)
```

Out[185]:

	FIPS	STFIPS	COFIPS	state	county	year	pop	white_pop	black_pop	asian_pop	india
0	-1.928755	-1.92287	-1.726894	0.458593	0.274188	-1.610567	0.616111	0.527443	1.420259	0.566839	0.0
1	-1.928755	-1.92287	-1.726894	0.458593	0.274188	-1.230735	0.629712	0.537918	1.432322	0.625005	0.0
2	-1.928755	-1.92287	-1.726894	0.458593	0.274188	-0.860716	0.624072	0.527914	1.436322	0.693883	0.0
3	-1.928755	-1.92287	-1.726894	0.458593	0.274188	-0.499397	0.619361	0.518138	1.438423	0.709458	0.

| 4 | -1.928755 | -1.92287 | -1.726894 | 0.458593 | 0.274188 | -0.145869 | 0.622808 | 0.518191 | 1.445458 | 0.729805 | 0. |

```
In [51]: x = df_pop.drop('pop', axis=1)
         y = df_pop['pop']
```

```
In [52]: sk = SelectKBest(score_func=f_regression,k=14)
         x_new = sk.fit_transform(x,y)
```

```
In [53]: #Get selected feture names and scores
         selected_features = x.columns[sk.get_support()]
         features_scores = pd.DataFrame({'feature':x.columns,'Score':sk.scores_}).sort_values(by=

         print_title('Selected Features:')
         print_section(pd.DataFrame(list(selected_features)))
         print_title("\nFeature Scores:")
         print_section(features_scores)
```

```
         ------------------------------------------------------------
         Selected Features:
         ------------------------------------------------------------
                          0
         0            FIPS
         1          STFIPS
         2          COFIPS
         3           state
         4          county
         5            year
         6       white_pop
         7       black_pop
         8       asian_pop
         9      indian_pop
         10    pacific_pop
         11        two_pop
         12   not_hisp_pop
         13       hisp_pop
         ------------------------------------------------------------


         ------------------------------------------------------------

         Feature Scores:
         ------------------------------------------------------------
                 feature          Score
         12   not_hisp_pop   1.540729e+06
         6       white_pop   1.164183e+06
         11        two_pop   2.898884e+05
         8       asian_pop   1.947178e+05
         10    pacific_pop   8.474241e+04
         13       hisp_pop   5.658111e+04
         9      indian_pop   3.518288e+04
         7       black_pop   3.100110e+04
         4          county   2.508289e+04
         3           state   1.200154e+04
         2          COFIPS   3.584230e+02
         0            FIPS   5.526506e-01
         1          STFIPS   4.202460e-01
         5            year   8.231369e-03
         ------------------------------------------------------------
```

```
In [54]: x_select=x[selected_features]
```

```
In [55]: x_select.columns
```

```
Out[55]: Index(['FIPS', 'STFIPS', 'COFIPS', 'state', 'county', 'year', 'white_pop',
                'black_pop', 'asian_pop', 'indian_pop', 'pacific_pop', 'two_pop',
```

```
                    'not_hisp_pop', 'hisp_pop'],
              dtype='object')
```

Feature Scaling

In [57]:
```python
# importing standardScaler from library
from sklearn.preprocessing import StandardScaler
```

In [58]:
```python
scaler = StandardScaler()
x_scaled = scaler.fit_transform(x_select)
```

Split data into Training and Testing Sets

In [60]:
```python
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
from sklearn.linear_model import LinearRegression
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import RandomizedSearchCV
from sklearn.svm import SVR
```

In [61]:
```python
# Split data into training and testing sets.
x_train, x_test, y_train, y_test = train_test_split(x_scaled, y, test_size=0.2, random_s
```

In [62]:
```python
print(f"""
shape x_train: {x_train.shape}
shape x_test: {x_test.shape}
shape y_train: {y_train.shape}
shape y_test: {y_test.shape}""")
```
```
shape x_train: (25128, 14)
shape x_test: (6282, 14)
shape y_train: (25128,)
shape y_test: (6282,)
```

ML Model

In [64]:
```python
# Creating a dictionary named models for model selection
models = {
    '1.linear Regression':LinearRegression(),
    '2.Dicision Tree Regression':DecisionTreeRegressor(),
    '3.Random Forest Regressor':RandomForestRegressor(),
    '4.Gradient Boosting Regressor':GradientBoostingRegressor(),
    '5.Support Vector Regressor':SVR()
}
```

In [65]:
```python
# This for loop syntax fit the x train and y train for each moledl in modles dictionary
result = {}
for model_name, model in models.items():
    model.fit(x_train,y_train)
    y_pred = model.predict(x_test)
```

```
        mae = mean_absolute_error(y_test,y_pred)
        mse = mean_squared_error(y_test,y_pred)
        rmse = np.sqrt(mse)
        r2 = r2_score(y_test,y_pred)
        result[model_name] = {'mae':mae, 'mse':mse,'rmse':rmse,'r2':r2}
```

Model Evaluation

In [67]:
```python
result_df = pd.DataFrame(result).T
print_title('Score details for variuse modles')
print_section(result_df)
```

```
------------------------------------------------------------
Score details for variuse modles
------------------------------------------------------------
                                    mae       mse      rmse        r2
1.linear Regression            0.041690  0.005840  0.076422  0.994167
2.Dicision Tree Regression     0.004423  0.000273  0.016536  0.999727
3.Random Forest Regressor      0.002218  0.000043  0.006553  0.999957
4.Gradient Boosting Regressor  0.017997  0.000744  0.027278  0.999257
5.Support Vector Regressor     0.044232  0.002742  0.052367  0.997261
------------------------------------------------------------
```

Hyperparameter turning

In [69]:
```python
rfg = RandomForestRegressor(random_state=42,
                            n_estimators=50,
                            max_depth=10,
                            min_samples_split=5,
                            min_samples_leaf=2,
                            max_features='sqrt',
                            ccp_alpha=0.1,
                            n_jobs=-1)
```

In [70]:
```python
param_grid = {
    'n_estimators':[50,100,200],
    'max_depth':[10,20,None],
    'min_samples_split':[2,5,10],
    'min_samples_leaf':[1,2,4]
}
```

In [71]:
```python
RandomGrid = RandomizedSearchCV(estimator = rfg,
                                param_distributions=param_grid,
                                cv=10,
                                scoring = 'r2',
                                n_jobs=-1,
                                verbose=2)
```

In [72]:
```python
RandomGrid.fit(x_train,y_train)
```

```
Fitting 10 folds for each of 10 candidates, totalling 100 fits
```

Out[72]:

```
   ▸          RandomizedSearchCV
                                        ⓘ ⓘ

   ▸            best_estimator_:
                RandomForestRegressor
        ▸ RandomForestRegressor  ⓘ
```

```
In [73]:  print("Best Parameters:",RandomGrid.best_params_)
          print("best R2 Score:",RandomGrid.best_score_)

          Best Parameters: {'n_estimators': 50, 'min_samples_split': 5, 'min_samples_leaf': 2, 'ma
          x_depth': None}
          best R2 Score: 0.8913701985915374
```

Save the Model

Best Model

```
In [140…   #creating best model variable using best estimator
           best_model = RandomGrid.best_estimator_
           best_model
```

```
Out[140]:             RandomForestRegressor                    ⓘ ⓘ

           RandomForestRegressor(ccp_alpha=0.1, max_features='sqrt', min_samples_leaf=2,
                                 min_samples_split=5, n_estimators=50, n_jobs=-1,
                                 random_state=42)
```

```
In [142…   # applaing paramiter to RandomForestRegrssor genarated from RandomGrid
           best_model = RandomForestRegressor(ccp_alpha=0.1,
                                              max_features='sqrt',
                                              min_samples_leaf=2,
                                              min_samples_split=5, n_estimators=50,
                                              n_jobs=-1,
                                              random_state=42)
```

```
In [144…   #fitting best model to traing data
           best_model.fit(x_train,y_train)
```

```
Out[144]:             RandomForestRegressor                    ⓘ ⓘ

           RandomForestRegressor(ccp_alpha=0.1, max_features='sqrt', min_samples_leaf=2,
                                 min_samples_split=5, n_estimators=50, n_jobs=-1,
                                 random_state=42)
```

```
In [146…   # genarating predictin from traing best model
           y_pred = best_model.predict(x_test)
```

```
In [150…   MAE = mean_absolute_error(y_test,y_pred)
           MSE = mean_squared_error(y_test,y_pred)
           RMSE = np.sqrt(MSE)
           R2 = r2_score(y_test,y_pred)
           print(f"MAE:{MAE}\nMSE:{MSE}\nRMSE:{RMSE}\nR2:{R2}")
```

```
MAE:0.2815571913338196
MSE:0.11309725190586278
RMSE:0.336293486551272
R2:0.8870478371921949
```

Model saving steps

In [167...
```python
# inporting library to save the model
import joblib
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer
```

In [165...
```python
pipeline  = Pipeline([
    ('imputer',SimpleImputer(strategy = 'mean')),('model',best_model)
])
```

In [157...
```python
pipeline.fit(x_train,y_train)
```

Out[157]:
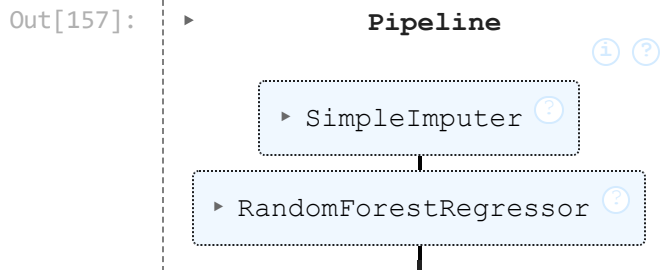
In [159...
```python
y_pred = pipeline.predict(x_test)
```

In [163...
```python
# saving the tested model as pop-data.joblib
joblib.dump(pipeline,'pop_data.joblib')
```

Out[163]:
```
['pop_data.joblib']
```

Test with Unseen Data

In [223...
```python
# Select 12 random indices
random_indices = np.random.choice(x.shape[0], 12, replace=False)
# Use NumPy indexing directly
x_random = x.iloc[random_indices]
y_random = y[random_indices]
```

In [229...
```python
# Apply the SelectKBest transformation to the random selection of features
x_random_selected = sk.transform(x_random)
```

In [231...
```python
prdiction = load_pipeline.predict(x_random_selected)
```

In [233...
```python
prdiction
```

Out[233]:
```
array([-0.76455391, -0.63569914,  1.1087055 , -0.30099588,  1.19104726,
       -0.47352545,  1.19104726, -0.71486351, -0.76455391, -0.52833682,
        1.19104726,  0.62945969])
```

In [257...
```python
result = pd.DataFrame(
    {
        "county":df.loc[x_random.index,'county'],
        "pop":y_random.values,
```

```
            "Predicted pop":prdiction
        }
    )
```

Final Result

```
In [261…  print(f'''Final Result
          {result}
          ''')
```

```
Final Result
               county        pop   Predicted pop
2754             Kiowa  -1.452569      -0.764554
2995          Saguache  -1.139307      -0.635699
17629     Carson City   0.643814       1.108705
21308            Adair  -0.340800      -0.300996
25826           Denton   1.604283       1.191047
28324         Buchanan  -0.303015      -0.473525
4184              Cobb   1.604283       1.191047
6673             Piatt  -0.595994      -0.714864
17513           Eureka  -1.407205      -0.764554
24552         Grainger  -0.321767      -0.528337
7598            Porter   1.604283       1.191047
13429           Isanti   0.280272       0.629460
```